

MINOR PROJECT

TASK 1 - Exploratory Data Analysis

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv("/content/HepatitisCdata.csv")
df.head(5)
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT	
0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0	
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17	4.80	74.0	15.6	76.5	
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84	5.20	86.0	33.2	79.3	
3	4	0=Blood Donor	32	m	43.2	52.0	30.6	22.6	18.9	7.33	4.74	80.0	33.8	75.7	
4	5	0=Blood Donor	32	m	39.2	74.1	32.6	24.8	9.6	9.15	4.32	76.0	29.9	68.7	

-----Question 1----->

QUESTION 1

- Are there any missing values in the dataset?
- If so, how will you handle them?
- What is the distribution of different diagnosis categories in the dataset?

```
#Are there any missing values in the dataset?
df.isnull()
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT	
0	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
...	
610	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
611	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
612	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
613	False	False	False	False	False	True	False	False	False	False	False	False	False	False	
614	False	False	False	False	False	True	False	False	False	False	False	False	False	False	

615 rows x 14 columns

```
df.isnull().sum()

Unnamed: 0      0
Category        0
Age             0
Sex             0
ALB             1
ALP            18
ALT             1
```

```

AST      0
BIL      0
CHE      0
CHOL     10
CREA     0
GGT      0
PROT     1
dtype: int64

```

therefore, from the above analysis we get the column names and their respective count of null values.

In order to handle these we can fill in the missing values as "NA"

```
df=df.fillna(value="NA")
```

```
df.isnull()
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT
0	False	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False
...
610	False	False	False	False	False	False	False	False	False	False	False	False	False	False
611	False	False	False	False	False	False	False	False	False	False	False	False	False	False
612	False	False	False	False	False	False	False	False	False	False	False	False	False	False
613	False	False	False	False	False	False	False	False	False	False	False	False	False	False
614	False	False	False	False	False	False	False	False	False	False	False	False	False	False

615 rows × 14 columns

```
df.isnull().sum()
```

```

Unnamed: 0    0
Category      0
Age           0
Sex           0
ALB           0
ALP           0
ALT           0
AST           0
BIL           0
CHE           0
CHOL          0
CREA          0
GGT           0
PROT          0
dtype: int64

```

#What is the distribution of different diagnosis categories in the dataset?

```

data=df['Category'].value_counts()
data

```

```

0=Blood Donor      533
3=Cirrhosis        30
1=Hepatitis        24
2=Fibrosis         21
0s=suspect Blood Donor  7
Name: Category, dtype: int64

```

Summarizing your analysis and observation

<-----Question 2----->

Perform feature encoding or transformation on categorical variables (such as sex) in the dataset. Which encoding technique would be most suitable?

```
df.head()
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT	
0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0	
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17	4.8	74.0	15.6	76.5	
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84	5.2	86.0	33.2	79.3	
3	4	0=Blood Donor	32	m	43.2	52.0	30.6	22.6	18.9	7.33	4.74	80.0	33.8	75.7	
4	5	0=Blood Donor	32	m	39.2	74.1	32.6	24.8	9.6	9.15	4.32	76.0	29.9	68.7	

```
encoded_data = pd.get_dummies(df, columns=["Sex"])
print(encoded_data.tail(10))
```

	Unnamed: 0	Category	Age	ALB	ALP	ALT	AST	BIL	CHE	\
605	606	3=Cirrhosis	42	33.0	79.0	3.7	55.7	200.0	1.72	
606	607	3=Cirrhosis	49	33.0	190.7	1.2	36.3	7.0	6.92	
607	608	3=Cirrhosis	52	39.0	37.0	1.3	30.4	21.0	6.33	
608	609	3=Cirrhosis	58	34.0	46.4	15.0	150.0	8.0	6.26	
609	610	3=Cirrhosis	59	39.0	51.3	19.6	285.8	40.0	5.77	
610	611	3=Cirrhosis	62	32.0	416.6	5.9	110.3	50.0	5.57	
611	612	3=Cirrhosis	64	24.0	102.8	2.9	44.4	20.0	1.54	
612	613	3=Cirrhosis	64	29.0	87.3	3.5	99.0	48.0	1.66	
613	614	3=Cirrhosis	46	33.0	NA	39.0	62.0	20.0	3.56	
614	615	3=Cirrhosis	59	36.0	NA	100.0	80.0	12.0	9.07	

	CHOL	CREA	GGT	PROT	Sex_f	Sex_m
605	5.16	89.1	146.3	69.9	1	0
606	3.82	485.9	112.0	58.5	1	0
607	3.78	158.2	142.5	82.7	1	0
608	3.98	56.0	49.7	80.6	1	0
609	4.51	136.1	101.1	70.5	1	0
610	6.3	55.7	650.9	68.5	1	0
611	3.02	63.0	35.9	71.3	1	0
612	3.63	66.7	64.2	82.0	1	0
613	4.2	52.0	50.0	71.0	1	0
614	5.3	67.0	34.0	68.0	1	0

```
# Add more cells if required
```

Summarizing your analysis and observations

<-----Question 3----->

Analyze the distribution and range of values for each clinical and demographic feature in the dataset. Are there any outliers or extreme values?

```
df.head()
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT
0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17	4.8	74.0	15.6	76.5
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84	5.2	86.0	33.2	79.3
3	4	0=Blood Donor	32	m	43.2	52.0	30.6	22.6	18.9	7.33	4.74	80.0	33.8	75.7
4	5	0=Blood Donor	32	m	39.2	74.1	32.6	24.8	9.6	9.15	4.32	76.0	29.9	68.7

Summarizing your analysis and observation

```
df.shape
```

```
(615, 14)
```

```
df.describe()
```

	Unnamed: 0	Age	AST	BIL	CHE	CREA	GGT
count	615.000000	615.000000	615.000000	615.000000	615.000000	615.000000	615.000000
mean	308.000000	47.408130	34.786341	11.396748	8.196634	81.287805	39.533171
std	177.679487	10.055105	33.090690	19.673150	2.205657	49.756166	54.661071
min	1.000000	19.000000	10.600000	0.800000	1.420000	8.000000	4.500000
25%	154.500000	39.000000	21.600000	5.300000	6.935000	67.000000	15.700000
50%	308.000000	47.000000	25.900000	7.300000	8.260000	77.000000	23.300000
75%	461.500000	54.000000	32.900000	11.200000	9.590000	88.000000	40.200000
max	615.000000	77.000000	324.000000	254.000000	16.410000	1079.100000	650.900000

```
df.columns
```

```
Index(['Unnamed: 0', 'Category', 'Age', 'Sex', 'ALB', 'ALP', 'ALT', 'AST',
      'BIL', 'CHE', 'CHOL', 'CREA', 'GGT', 'PROT'],
      dtype='object')
```

```
features=["Age", "ALB", "ALP", "AST", "BIL", "CHE", "CHOL", "CREA", "GGT", "PROT"]
```

```
for feature in features:
```

```
    print(feature)
```

```
    plt.figure(figsize=(5,5))
```

```
    plt.hist(pd.to_numeric(df[feature], errors="coerce"), bins=20)
```

```
    plt.xlabel(feature.capitalize())
```

```
    plt.ylabel("Frequency")
```

```
    plt.title(f"Distribution of {feature.capitalize()}")
```

```
    plt.show()
```

```
min=df[feature].min
```


```
max=df[feature].max
```

```
min_value = pd.to_numeric(df[feature], errors="coerce").min()
```

```
max_value = pd.to_numeric(df[feature], errors="coerce").max()
```


```
print(f"Range of values for {feature.capitalize()}: {min_value} - {max_value}")
```

Age



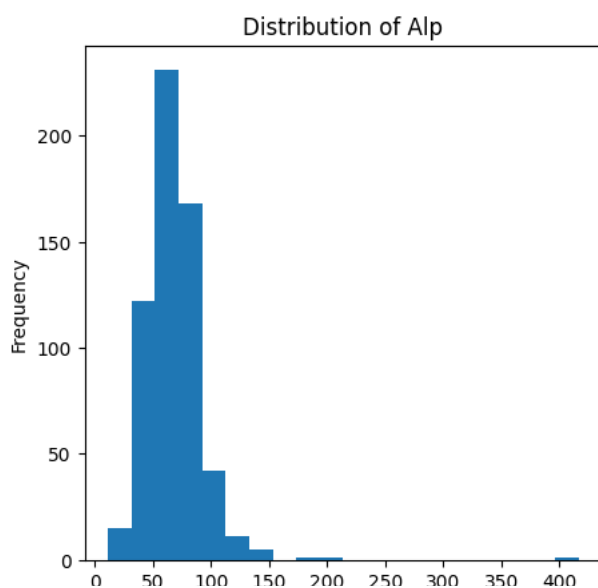
Range of values for Age: 19 - 77

ALB




Range of values for Alb: 14.9 - 82.2


ALP



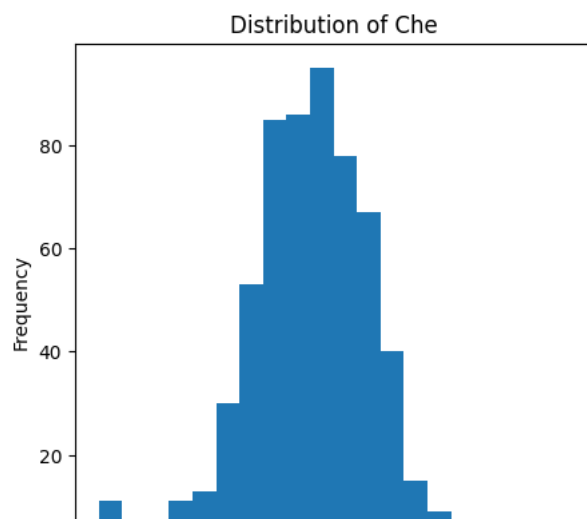
Alp
Range of values for Alp: 11.3 - 416.6
AST




Range of values for Ast: 10.6 - 324.0
BIL




Range of values for Bil: 0.8 - 254.0
CHE

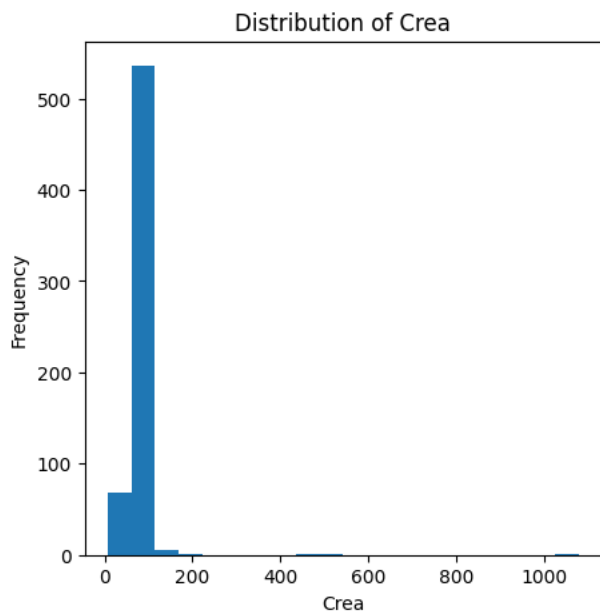





Range of values for Che: 1.42 - 16.41
CHOL




Range of values for Chol: 1.43 - 9.67
CREA

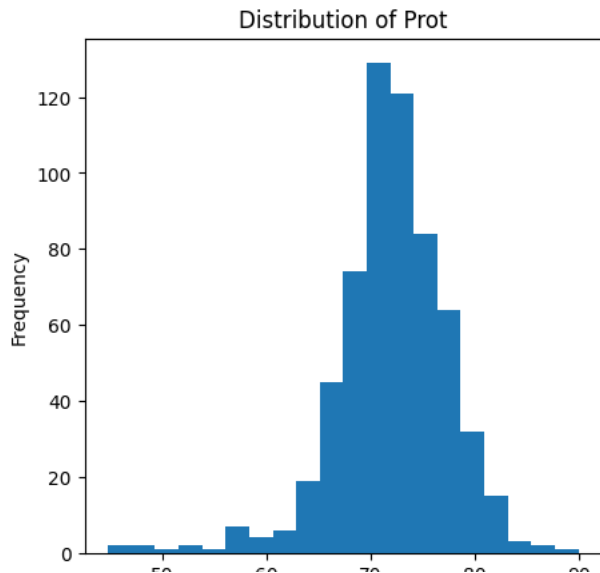


Range of values for Crea: 8.0 - 1079.1
GGT





Range of values for Ggt: 4.5 - 650.9
PROT



```
from scipy import stats
```

```
# Calculate z-scores for the "price" variable
z_scores = stats.zscore(pd.to_numeric(df["GGT"],errors="coerce"))
```

```
# Set a threshold for outliers (e.g., z-score > 3 or < -3)
threshold = 3
```

```
# Identify the outliers
outliers = df[abs(z_scores) > threshold]
```

```
# Remove the outliers from the dataset
clean_data = df[abs(z_scores) <= threshold]
```

```
# Display the outliers and cleaned dataset
print("Outliers:")
print(outliers)
print("\nCleaned Dataset:")
print(clean_data.head())
```

Outliers:

Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	\
205	0=Blood Donor	50	m	42.2	145.0	27.5	37.9	
533	0s=suspect Blood Donor	47	m	22.5	124.0	79.5	46.7	
538	0s=suspect Blood Donor	74	m	20.3	84.0	22.8	43.0	
539	0s=suspect Blood Donor	59	f	19.3	208.2	325.3	146.6	
558	1=Hepatitis	56	m	37.0	114.0	27.8	324.0	
559	1=Hepatitis	58	m	43.0	99.1	12.2	63.2	
593	3=Cirrhosis	51	m	39.0	66.0	29.6	185.0	
598	3=Cirrhosis	58	m	31.0	143.1	7.0	181.8	
602	3=Cirrhosis	61	m	39.0	102.9	27.3	143.2	
610	3=Cirrhosis	62	f	32.0	416.6	5.9	110.3	

	BIL	CHE	CHOL	CREA	GGT	PROT
205	4.5	13.71	8.8	103.0	239.0	73.1
533	2.3	6.83	4.3	170.0	345.6	58.6
538	5.7	4.91	3.19	52.0	218.3	47.8
539	6.9	5.33	4.72	32.0	295.6	53.1
558	67.0	5.75	3.09	97.7	392.2	77.3
559	13.0	5.95	6.15	147.3	491.0	65.6
593	19.0	2.00	3.6	58.3	399.5	79.4
598	58.0	3.29	3.92	66.4	273.7	78.1
602	15.0	5.38	4.88	72.3	400.3	73.4
610	50.0	5.57	6.3	55.7	650.9	68.5

Cleaned Dataset:

Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	\
------------	----------	-----	-----	-----	-----	-----	-----	-----	-----	---

0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84
3	4	0=Blood Donor	32	m	43.2	52.0	30.6	22.6	18.9	7.33
4	5	0=Blood Donor	32	m	39.2	74.1	32.6	24.8	9.6	9.15

	CHOL	CREA	GGT	PROT
0	3.23	106.0	12.1	69.0
1	4.8	74.0	15.6	76.5
2	5.2	86.0	33.2	79.3
3	4.74	80.0	33.8	75.7
4	4.32	76.0	29.9	68.7

<-----Question 4----->

```
#-----code

# Add more cells if required
```

Summarizing your analysis and observations

<-----Question 5----->

Can you explore the correlations between features and the diagnosis of hepatitis C? Are there any strong correlations or dependencies?

```
df.info

<bound method DataFrame.info of
0      1  0=Blood Donor  32  m  38.5  52.5   7.7  22.1   7.5   6.93
1      2  0=Blood Donor  32  m  38.5  70.3  18.0  24.7   3.9  11.17
2      3  0=Blood Donor  32  m  46.9  74.7  36.2  52.6   6.1   8.84
3      4  0=Blood Donor  32  m  43.2  52.0  30.6  22.6  18.9   7.33
4      5  0=Blood Donor  32  m  39.2  74.1  32.6  24.8   9.6   9.15
..    ...
610   611   3=Cirrhosis  62  f  32.0  416.6   5.9 110.3  50.0
611   612   3=Cirrhosis  64  f  24.0  102.8   2.9  44.4  20.0
612   613   3=Cirrhosis  64  f  29.0   87.3   3.5  99.0  48.0
613   614   3=Cirrhosis  46  f  33.0    NA  39.0  62.0  20.0
614   615   3=Cirrhosis  59  f  36.0    NA 100.0  80.0  12.0

      CHE  CHOL  CREA  GGT  PROT
0      6.93  3.23 106.0 12.1 69.0
1     11.17  4.8  74.0 15.6 76.5
2      8.84  5.2  86.0 33.2 79.3
3      7.33  4.74  80.0 33.8 75.7
4      9.15  4.32  76.0 29.9 68.7
..    ...
610   5.57  6.3  55.7 650.9 68.5
611   1.54  3.02  63.0  35.9 71.3
612   1.66  3.63  66.7  64.2 82.0
613   3.56  4.2  52.0  50.0 71.0
614   9.07  5.3  67.0  34.0 68.0

[615 rows x 14 columns]>
```

```
df.head()
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT
0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17	4.8	74.0	15.6	76.5
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84	5.2	86.0	33.2	79.3
3	4	0=Blood Donor	32	m	43.2	52.0	30.6	22.6	18.9	7.33	4.74	80.0	33.8	75.7
4	5	0=Blood Donor	32	m	39.2	74.1	32.6	24.8	9.6	9.15	4.32	76.0	29.9	68.7

```
selected_columns = ["Age", "ALB", "ALP", "ALT", "AST", "BIL", "CHE", "CHOL", "CREA", "GGT", "PROT"]
subset_df = df[selected_columns]
subset_df.head()
```

	Age	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT
0	32	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0
1	32	38.5	70.3	18.0	24.7	3.9	11.17	4.8	74.0	15.6	76.5
2	32	46.9	74.7	36.2	52.6	6.1	8.84	5.2	86.0	33.2	79.3
3	32	43.2	52.0	30.6	22.6	18.9	7.33	4.74	80.0	33.8	75.7
4	32	39.2	74.1	32.6	24.8	9.6	9.15	4.32	76.0	29.9	68.7

```
correlation_matrix = subset_df.corr()
correlation_matrix
```


<ipython-input-33-bda98bf6fb1c>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False, meaning non-numeric data will be included in the correlation calculation.

```
correlation_matrix = subset_df.corr()
```

	Age	AST	BIL	CHE	CREA	GGT
Age	1.000000	0.088666	0.032492	-0.075093	-0.022296	0.153087
AST	0.088666	1.000000	0.312231	-0.208536	-0.021387	0.491263
BIL	0.032492	0.312231	1.000000	-0.333172	0.031224	0.217024
CHE	-0.075093	-0.208536	-0.333172	1.000000	-0.011157	-0.110345
CREA	-0.022296	-0.021387	0.031224	-0.011157	1.000000	0.121003
GGT	0.153087	0.491263	0.217024	-0.110345	0.121003	1.000000

```
import seaborn as sns
```

```
# Create a heatmap of the correlation matrix
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm")
```



Summarizing your analysis and observation

▼ TASK 2 - Classification/Regression

Perform following steps on the same dataset which you used for EDA.

- Data Preprocessing (as per requirement)
- Feature Engineering
- Split dataset in train-test (80:20 ratio)
- Model selection
- Model training
- Model evaluation
- Fine-tune the Model
- Make predictions

Summarize your model's performance by evaluation metrices

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
data=pd.read_csv("/content/HepatitisCdata.csv")
data.head(5)
```

	Unnamed: 0	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	CHOL	CREA	GGT	PROT
0	1	0=Blood Donor	32	m	38.5	52.5	7.7	22.1	7.5	6.93	3.23	106.0	12.1	69.0
1	2	0=Blood Donor	32	m	38.5	70.3	18.0	24.7	3.9	11.17	4.80	74.0	15.6	76.5
2	3	0=Blood Donor	32	m	46.9	74.7	36.2	52.6	6.1	8.84	5.20	86.0	33.2	79.3

```
data.isnull().sum()
```

```
Unnamed: 0      0
Category        0
Age             0
Sex             0
ALB             1
ALP            18
ALT             1
AST             0
BIL             0
CHE             0
CHOL           10
CREA            0
GGT             0
PROT            1
dtype: int64
```

```
from sklearn.model_selection import train_test_split
```

```
# Separate the target variable from the features
```

```
target = data['Category']
```

```
features = data.drop('Category', axis=1)
```

```
# Split the dataset into train and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# Create the Random Forest Classifier
```

```

model = RandomForestClassifier()

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
import pandas as pd

# Load the dataset
data=pd.read_csv("/content/HepatitisCdata.csv")

# Convert 'Sex' column to numerical values
data = pd.get_dummies(data, columns=['Sex'], drop_first=True)

# Separate the target variable from the features
target = data['Category']
features = data.drop('Category', axis=1)

# Perform missing value imputation
imputer = SimpleImputer(strategy='mean')
imputed_features = pd.DataFrame(imputer.fit_transform(features), columns=features.columns)

# Split the dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(imputed_features, target, test_size=0.2, random_state=42)

# Create the Random Forest Classifier
model = RandomForestClassifier()

# Train the model
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
report = classification_report(y_test, y_pred)
print(report)

```

	precision	recall	f1-score	support
0=Blood Donor	0.99	1.00	0.99	96
0s=suspect Blood Donor	1.00	0.33	0.50	3
1=Hepatitis	0.80	0.89	0.84	9
2=Fibrosis	0.62	0.83	0.71	6
3=Cirrhosis	1.00	0.78	0.88	9
accuracy			0.95	123
macro avg	0.88	0.77	0.79	123
weighted avg	0.96	0.95	0.95	123

```

from sklearn.metrics import classification_report

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
report = classification_report(y_test, y_pred)
print(report)

```

	precision	recall	f1-score	support
0=Blood Donor	0.99	1.00	0.99	96
0s=suspect Blood Donor	1.00	0.33	0.50	3
1=Hepatitis	0.80	0.89	0.84	9
2=Fibrosis	0.62	0.83	0.71	6
3=Cirrhosis	1.00	0.78	0.88	9
accuracy			0.95	123
macro avg	0.88	0.77	0.79	123
weighted avg	0.96	0.95	0.95	123

```
from sklearn.model_selection import GridSearchCV
```

```
# Define the hyperparameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10]
}
```

```
# Create the GridSearchCV object
grid_search = GridSearchCV(model, param_grid, cv=5)
```

```
# Perform grid search to find the best hyperparameters
grid_search.fit(X_train, y_train)
```

```
# Get the best model with tuned hyperparameters
best_model = grid_search.best_estimator_
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_split.py:700: UserWarning: The least populated class in y has only 4 members; please consider decreasing the number of classes.
warnings.warn(
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
```

```
# Display the performance summary
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
```

```
Accuracy: 0.9512195121951219
Precision: 0.9590269047020368
Recall: 0.9512195121951219
F1-score: 0.9491240026643128
```

✓ 0s completed at 11:45 PM

